## Car Prices

- The goal of this challenge is to prepare a dataset and apply some feature selection techniques that you have learned so far.
- ## We are dealing with a dataset about cars and we would like to predict whether a car is expensive or cheap.

```
In [1]:
          # Data manipulation
          import numpy as np
          import pandas as pd
          # Data visualisation
          import matplotlib.pyplot as plt
          import seaborn as sns
          # Checking whether a numerical feature has a normal distribution or not
          from statsmodels.graphics.gofplots import qqplot
 In [2]: url = "https://wagon-public-datasets.s3.amazonaws.com/Machine%20Learning%20Dataset
          ? Go ahead and load the CSV into a dataframe called df.
In [13]: df = pd.read_csv(url)
          df.head(2)
Out[13]:
                       enginelocation carwidth curbweight enginetype cylindernumber stroke peakrpn
          0
                                         64.1
                                                    2548
                                                                                      2.68
                                                                                               500
                               front
                                                               dohc
                                                                               four
                   std
                   std
                               front
                                         64.1
                                                    2548
                                                               dohc
                                                                               four
                                                                                      2.68
                                                                                               500
```

i The description of the dataset is available here. Make sure to refer to it throughout the exercise.

# (1) Duplicates

? Remove the duplicates from the dataset if there are any. ?

Overwite the dataframe df

```
In [14]: df= df.drop_duplicates().reset_index(drop = True)
```

## (2) Missing values

? Find the missing values and impute them either with strategy = "most frequent" (categorical variables) or strategy = "median" (numerical variables) ?

```
In [15]: round(df.isnull().sum().sort_values(ascending=False)/len(df),3)
```

 $\hbox{\it enginelocation}$ 0.052 Out[15]: carwidth 0.010 aspiration 0.000 curbweight 0.000 enginetype 0.000 cylindernumber 0.000 stroke 0.000 peakrpm 0.000 0.000 price dtype: float64

## carwidth

#### ▶ 💡 Hint

```
In [16]: df.carwidth.value_counts(dropna=False)
```

```
66.5
                   22
Out[16]:
          63.8
                   19
          65.4
                   15
          63.6
                    9
          68.4
                    9
          64
                    9
          64.4
                    9
          65.5
                    8
          65.2
                    7
          65.6
                    6
          64.2
                    6
          66.3
                    6
          67.2
                    6
          66.9
                    5
                    5
          67.9
                    4
                    4
          68.9
          71.7
                    3
          70.3
                    3
          65.7
                    3
                    3
          63.9
          64.8
                    3
          65
                    2
          67.7
                    2
                    2
          68.3
          71.4
                    2
          NaN
                    2
          66.6
                    1
          63.4
                    1
          72.3
                    1
          64.1
                    1
          68
                    1
          72
                    1
          70.5
                    1
          66.1
                    1
          70.6
                    1
          69.6
                    1
          61.8
                    1
          66
                    1
          64.6
                    1
                    1
          60.3
          70.9
                    1
          66.4
                    1
          68.8
```

Name: carwidth, dtype: int64

```
import numpy as np
from sklearn.impute import SimpleImputer

df = df.replace("*", np.nan) # Replace occurences of "*" by np.nan

carwidth_imputer = SimpleImputer(strategy="median") # Instanciate median imputer
    carwidth_imputer.fit(df[['carwidth']]) # Fit imputer to carwidth column
    df['carwidth'] = carwidth_imputer.transform(df[['carwidth']]) # Impute

df.head()
```

Out[17]:		aspiration	enginelocation	carwidth	curbweight	enginetype	cylindernumber	stroke	peakrpn
	0	std	front	64.1	2548	dohc	four	2.68	500
	1	std	front	65.5	2823	ohcv	six	3.47	500
	2	std	front	65.5	2337	ohc	four	3.40	550
	3	std	front	66.4	2824	ohc	five	3.40	550
	4	std	front	66.3	2507	ohc	five	3.40	550
4									<b>•</b>

### enginelocation

#### ▶ 💡 Hint

```
In [18]: print(df.enginelocation.unique())
    print(df.enginelocation.value_counts(dropna=False))

['front' nan 'rear']
    front 179
    NaN 10
    rear 2
```

Name: enginelocation, dtype: int64

In [20]: engine\_imputer = SimpleImputer(strategy="most\_frequent") # Instantiate most freque
engine\_imputer.fit(df[['enginelocation']]) # Fit imputer to enginelocation column
df['enginelocation'] = engine\_imputer.transform(df[['enginelocation']]) # Impute

df.head()

Out[20]:		aspiration	enginelocation	carwidth	curbweight	enginetype	cylindernumber	stroke	peakrpn
	0	std	front	64.1	2548	dohc	four	2.68	500
	1	std	front	65.5	2823	ohcv	six	3.47	500
	2	std	front	65.5	2337	ohc	four	3.40	550
	3	std	front	66.4	2824	ohc	five	3.40	550
	4	std	front	66.3	2507	ohc	five	3.40	550
4									<b>&gt;</b>

#### Test your code

```
In [21]: from nbresult import ChallengeResult
    result = ChallengeResult('missing_values',
```

```
dataset = df
result.write()
print(result.check())
platform linux -- Python 3.10.6, pytest-7.1.3, pluggy-1.0.0 -- /home/joharlewago
n/.pyenv/versions/lewagon/bin/python3
cachedir: .pytest cache
rootdir: /home/joharlewagon/code/UKVeteran/05-ML/02-Prepare-the-dataset/data-car-p
rices/tests
plugins: anyio-3.6.2, asyncio-0.19.0, typeguard-2.13.3
asyncio: mode=strict
collecting ... collected 2 items
test_missing_values.py::TestMissing_values::test_carwidth PASSED
test_missing_values.py::TestMissing_values::test_engine_location PASSED [100%]
You can commit your code:
git add tests/missing_values.pickle
git commit -m 'Completed missing values step'
git push origin master
```

# (3) Scaling the numerical features

```
In [22]: # As a reminder, some information about the dataframe
         df.info()
         <class 'pandas.core.frame.DataFrame'>
         RangeIndex: 191 entries, 0 to 190
         Data columns (total 9 columns):
          # Column
                             Non-Null Count Dtype
         --- -----
                             _____
             aspiration
                             191 non-null
                                            object
          1 enginelocation 191 non-null
                                            object
          2 carwidth
                           191 non-null
                                            float64
          3
            curbweight
                            191 non-null
                                            int64
                             191 non-null
            enginetype
                                            object
             cylindernumber 191 non-null
                                            object
          6
             stroke
                             191 non-null
                                            float64
                             191 non-null
             peakrpm
                                            int64
             price
                             191 non-null
                                             object
         dtypes: float64(2), int64(2), object(5)
         memory usage: 13.6+ KB
         # And here are the numerical features of the dataset we need to scale
In [23]:
         numerical_features = df.select_dtypes(exclude=['object']).columns
         numerical_features
         Index(['carwidth', 'curbweight', 'stroke', 'peakrpm'], dtype='object')
Out[23]:
```

#### ? Question: Scaling the numerical features ?

Investigate the numerical features for outliers and distribution, and apply the solutions below accordingly:

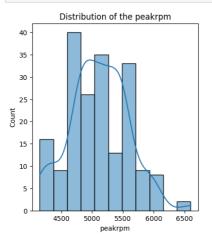
- Robust Scaler
- Standard Scaler

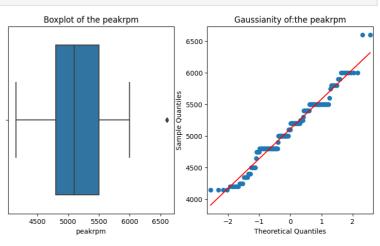
Replace the original columns with the transformed values.

# peakrpm , carwidth , & stroke

#### Hint

```
In [25]: variable = 'peakrpm'
fig, ax = plt.subplots(1,3,figsize=(15,5))
ax[0].set_title(f"Distribution of the {variable}")
sns.histplot(data = df, x = f"{variable}", kde=True, ax = ax[0])
ax[1].set_title(f"Boxplot of the {variable}")
sns.boxplot(data = df, x = f"{variable}", ax=ax[1])
ax[2].set_title(f"Gaussianity of:the {variable}")
qqplot(df[f"{variable}"],line='s',ax=ax[2]);
```





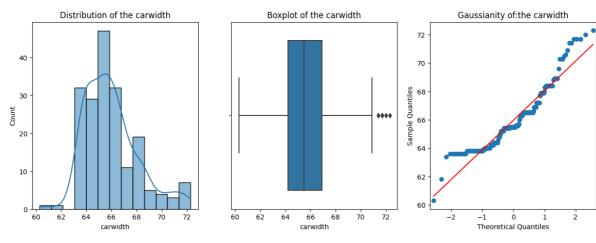
```
In [26]: variable = 'carwidth'

fig, ax = plt.subplots(1,3,figsize=(15,5))

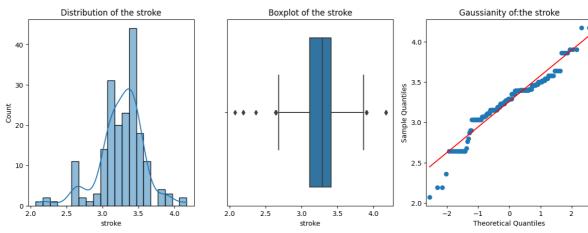
ax[0].set_title(f"Distribution of the {variable}")
sns.histplot(data = df, x = f"{variable}", kde=True, ax = ax[0])

ax[1].set_title(f"Boxplot of the {variable}")
sns.boxplot(data = df, x = f"{variable}", ax=ax[1])

ax[2].set_title(f"Gaussianity of:the {variable}")
qqplot(df[f"{variable}"],line='s',ax=ax[2]);
```



```
In [27]: variable = 'stroke'
fig, ax = plt.subplots(1,3,figsize=(15,5))
ax[0].set_title(f"Distribution of the {variable}")
sns.histplot(data = df, x = f"{variable}", kde=True, ax = ax[0])
ax[1].set_title(f"Boxplot of the {variable}")
sns.boxplot(data = df, x = f"{variable}", ax=ax[1])
ax[2].set_title(f"Gaussianity of:the {variable}")
qqplot(df[f"{variable}"],line='s',ax=ax[2]);
```



```
In [28]: from sklearn.preprocessing import RobustScaler

rb_scaler = RobustScaler()
df['peakrpm'],df['carwidth'],df['stroke'] = rb_scaler.fit_transform(df[['peakrpm',
df.head()
```

Out[28]:		aspiration	enginelocation	carwidth	curbweight	enginetype	cylindernumber	stroke	pea
	0	std	front	-0.518519	2548	dohc	four	-2.033333	-0.14
	1	std	front	0.000000	2823	ohcv	six	0.600000	-0.14
	2	std	front	0.000000	2337	ohc	four	0.366667	0.57
	3	std	front	0.333333	2824	ohc	five	0.366667	0.57
	4	std	front	0.296296	2507	ohc	five	0.366667	0.57

#### ▶ 🦞 Hint

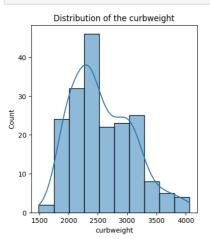
```
In [29]: variable = 'curbweight'

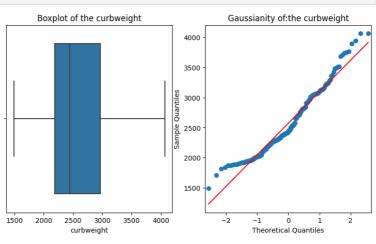
fig, ax = plt.subplots(1,3,figsize=(15,5))

ax[0].set_title(f"Distribution of the {variable}")
    sns.histplot(data = df, x = f"{variable}", kde=True, ax = ax[0])

ax[1].set_title(f"Boxplot of the {variable}")
    sns.boxplot(data = df, x = f"{variable}", ax=ax[1])

ax[2].set_title(f"Gaussianity of:the {variable}")
    qqplot(df[f"{variable}"],line='s',ax=ax[2]);
```





```
In [30]: from sklearn.preprocessing import StandardScaler

std_scaler = StandardScaler()
df['curbweight'] = std_scaler.fit_transform(df[['curbweight']])
df.head()
```

pea	stroke	cylindernumber	enginetype	curbweight	carwidth	enginelocation	aspiration	[30]:	Οι
-0.14	-2.033333	four	dohc	-0.048068	-0.518519	front	std	0	
-0.14	0.600000	six	ohcv	0.476395	0.000000	front	std	1	
0.57	0.366667	four	ohc	-0.450474	0.000000	front	std	2	
0.57	0.366667	five	ohc	0.478302	0.333333	front	std	3	
0.57	0.366667	five	ohc	-0.126260	0.296296	front	std	4	
									4

#### Test your code

# (4) Encoding the categorical features

- ? Question: encoding the categorical variables ?
- Investigate the features that require encoding, and apply the following techniques accordingly:
  - One-hot encoding
  - Manual ordinal encoding

In the Dataframe, replace the original features with their encoded version(s).

```
In [32]: print(f"The unique values of `aspiration` are {df.aspiration.unique()}") # Check us
print(f"The unique values of `enginelocation` are {df.enginelocation.unique()}") #
The unique values of `aspiration` are ['std' 'turbo']
The unique values of `enginelocation` are ['front' 'rear']
```

### aspiration & enginelocation

#### ► 🦞 Hint

```
In [33]: from sklearn.preprocessing import OneHotEncoder
binary_encoder = OneHotEncoder(sparse=False, drop='if_binary')
df['aspiration'], df['enginelocation'] = binary_encoder.fit_transform(df[['aspiration'], df['enginelocation'])
```

Out[33]:		aspiration	enginelocation	carwidth	curbweight	enginetype	cylindernumber	stroke	pea
	0	0.0	0.0	-0.518519	-0.048068	dohc	four	-2.033333	-0.14
	1	0.0	0.0	0.000000	0.476395	ohcv	six	0.600000	-0.14
	2	0.0	0.0	0.000000	-0.450474	ohc	four	0.366667	0.57
	3	0.0	0.0	0.333333	0.478302	ohc	five	0.366667	0.57
	4	0.0	0.0	0.296296	-0.126260	ohc	five	0.366667	0.57
4									

## enginetype

#### 🕨 💡 Hint

```
In [34]: df.enginetype.unique()
Out[34]: array(['dohc', 'ohcv', 'ohc', 'l', 'rotor', 'ohcf', 'dohcv'], dtype=object)
In [35]: df.shape
Out[35]: (191, 9)
```

```
from sklearn.preprocessing import OneHotEncoder
In [36]:
          # Instantiate a OneHotEncoder for the categorical feature EngineType
          ohe = OneHotEncoder(sparse=False)
          # Fitting it
          ohe.fit(df[['enginetype']])
          # Showing the categories detected by the encoder
          display(ohe.categories_)
          # Since Sklearn 1.1, we can retrieve the names of the generated columns
          display(ohe.get_feature_names_out())
          # Let's encode EngineType
          enginetype_encoded = ohe.transform(df[['enginetype']])
          # Now we store the encoded values in the dataframe
          df[ohe.get_feature_names_out()] = enginetype_encoded
          # We can get rid of the original column EngineType now
          df.drop(columns='enginetype', inplace = True)
          # And show df
          df
          [array(['dohc', 'dohcv', 'l', 'ohc', 'ohcf', 'ohcv', 'rotor'], dtype=object)]
          'enginetype_rotor'], dtype=object)
               aspiration enginelocation carwidth curbweight cylindernumber
Out[36]:
                                                                              stroke
                                                                                     peakrpm
            0
                     0.0
                                       -0.518519
                                                   -0.048068
                                                                           -2.033333
                                   0.0
                                                                      four
                                                                                     -0.142857
                                                                                              expe
            1
                     0.0
                                   0.0
                                        0.000000
                                                   0.476395
                                                                            0.600000
                                                                                    -0.142857 expe
                                                                       six
            2
                     0.0
                                   0.0
                                        0.000000
                                                   -0.450474
                                                                      four
                                                                            0.366667
                                                                                     0.571429 expe
                                        0.333333
                                                                            0.366667
            3
                     0.0
                                   0.0
                                                   0.478302
                                                                       five
                                                                                     0.571429
                                                                                              expe
            4
                     0.0
                                   0.0
                                        0.296296
                                                   -0.126260
                                                                       five
                                                                            0.366667
                                                                                     0.571429
                                                                                              expe
          186
                     0.0
                                        1.259259
                                   0.0
                                                   0.722416
                                                                      four
                                                                           -0.466667
                                                                                     0.428571
                                                                                              expe
          187
                     1.0
                                   0.0
                                        1.222222
                                                   0.907408
                                                                      four
                                                                           -0.466667
                                                                                     0.285714
                                                                                              expe
          188
                     0.0
                                        1.259259
                                   0.0
                                                   0.836844
                                                                           -1.400000
                                                                                     0.571429
                                                                                              expe
          189
                     1.0
                                   0.0
                                        1.259259
                                                   1.227807
                                                                            0.366667
                                                                                     -0.428571
                                                                                              expe
          190
                     1.0
                                   0.0
                                        1.259259
                                                   0.932201
                                                                           -0.466667
                                                                                     0.428571 expe
                                                                      four
         191 rows × 15 columns
```

## cylindernumber


Hint

In [37]: df.cylindernumber.unique()

Out[38]:		aspiration	enginelocation	carwidth	curbweight	cylindernumber	stroke	peakrpm	pr
	0	0.0	0.0	-0.518519	-0.048068	4	-2.033333	-0.142857	expens
	1	0.0	0.0	0.000000	0.476395	6	0.600000	-0.142857	expens
	2	0.0	0.0	0.000000	-0.450474	4	0.366667	0.571429	expens
	3	0.0	0.0	0.333333	0.478302	5	0.366667	0.571429	expens
	4	0.0	0.0	0.296296	-0.126260	5	0.366667	0.571429	expens
4				_					

? Now that you've made cylindernumber into a numeric feature between 2 and 12, you need to scale it ?

#### ► P Hint



```
In [42]: from sklearn.preprocessing import MinMaxScaler, RobustScaler
          mm_scaler = MinMaxScaler()
          mm = pd.DataFrame(mm_scaler.fit_transform(df[['cylindernumber']]))
          rb_scaler = RobustScaler()
          rb = pd.DataFrame(rb_scaler.fit_transform(df[['cylindernumber']]))
         mm.value_counts()
In [43]:
         0.2
                 147
Out[43]:
         0.4
                 23
         0.3
                  11
         0.6
         0.0
                   3
                   1
         0.1
         1.0
                   1
         dtype: int64
         rb.value_counts()
In [44]:
                 147
          0.0
Out[44]:
          2.0
                   23
          1.0
                   11
          4.0
                    5
          -2.0
                    3
          -1.0
                    1
          8.0
                    1
         dtype: int64
In [45]: from sklearn.preprocessing import RobustScaler
          rb_scaler = RobustScaler()
          df['cylindernumber'] = rb_scaler.fit_transform(df[['cylindernumber']])
          df.head()
```

Out[45]:		aspiration	enginelocation	carwidth	curbweight	cylindernumber	stroke	peakrpm	ıq
	0	0.0	0.0	-0.518519	-0.048068	0.0	-2.033333	-0.142857	expens
	1	0.0	0.0	0.000000	0.476395	2.0	0.600000	-0.142857	expens
	2	0.0	0.0	0.000000	-0.450474	0.0	0.366667	0.571429	expens
	3	0.0	0.0	0.333333	0.478302	1.0	0.366667	0.571429	expens
	4	0.0	0.0	0.296296	-0.126260	1.0	0.366667	0.571429	expens
4									<b>•</b>

▶ Here is a screenshot of how your dataframe shoud look like after scaling and encoding

## price

- Encode the target price.
- ► 🦞 Hint

```
In [46]: from sklearn.preprocessing import LabelEncoder

df['price'] = LabelEncoder().fit_transform(df['price'])
    df.head()
```

```
aspiration
                          enginelocation
                                           carwidth curbweight cylindernumber
Out[46]:
                                                                                        stroke
                                                                                                peakrpm price
           0
                                           -0.518519
                                                        -0.048068
                                                                                     -2.033333
                      0.0
                                      0.0
                                                                                0.0
                                                                                                -0.142857
                                                                                                               1
                      0.0
                                      0.0
                                            0.000000
                                                         0.476395
                                                                                2.0
                                                                                      0.600000
                                                                                                -0.142857
           2
                                            0.000000
                                                        -0.450474
                                                                                      0.366667
                      0.0
                                      0.0
                                                                                0.0
                                                                                                 0.571429
                                                                                                               1
            3
                      0.0
                                            0.333333
                                                         0.478302
                                                                                      0.366667
                                                                                                 0.571429
                                      0.0
            4
                      0.0
                                                        -0.126260
                                                                                      0.366667
                                                                                                 0.571429
                                                                                                               1
                                      0.0
                                            0.296296
                                                                                1.0
```

#### Test your code

```
================= test session starts ======================
platform linux -- Python 3.10.6, pytest-7.1.3, pluggy-1.0.0 -- /home/joharlewago
n/.pyenv/versions/lewagon/bin/python3
cachedir: .pytest_cache
rootdir: /home/joharlewagon/code/UKVeteran/05-ML/02-Prepare-the-dataset/data-car-p
rices/tests
plugins: anyio-3.6.2, asyncio-0.19.0, typeguard-2.13.3
asyncio: mode=strict
collecting ... collected 4 items
test encoding.py::TestEncoding::test_aspiration PASSED
                                                                 [ 25%]
test_encoding.py::TestEncoding::test_enginelocation PASSED
                                                                 [ 50%]
test_encoding.py::TestEncoding::test_enginetype PASSED
                                                                 [ 75%]
test_encoding.py::TestEncoding::test_price PASSED
                                                                 [100%]
You can commit your code:
git add tests/encoding.pickle
git commit -m 'Completed encoding step'
git push origin master
```

# (5) Base Modelling

nterior The dataset has been preprocessed and is now ready to be fitted to a model.

#### ? Question: a first attempt to evaluate a classification model ?

Cross-validate a LogisticRegression on this preprocessed dataset and save its score under a variable named base model score.

```
In [48]: from sklearn.linear_model import LogisticRegression
    from sklearn.model_selection import cross_val_score

X = df.drop(columns=['price'])
y = df['price']

model = LogisticRegression()

scores = cross_val_score(model, X, y, cv=10)
base_model_score = scores.mean()

base_model_score
```

#### Test your code

0.8797368421052632

Out[48]:

```
result.write()
print(result.check())
platform linux -- Python 3.10.6, pytest-7.1.3, pluggy-1.0.0 -- /home/joharlewago
n/.pyenv/versions/lewagon/bin/python3
cachedir: .pytest_cache
rootdir: /home/joharlewagon/code/UKVeteran/05-ML/02-Prepare-the-dataset/data-car-p
rices/tests
plugins: anyio-3.6.2, asyncio-0.19.0, typeguard-2.13.3
asyncio: mode=strict
collecting ... collected 1 item
test_base_model.py::TestBase_model::test_base_model_score PASSED
                                                         [100%]
You can commit your code:
git add tests/base_model.pickle
git commit -m 'Completed base_model step'
git push origin master
```

## (6) Feature Selection (with *Permutation Importance*)

- A powerful way to detect whether a feature is relevant or not to predict a target is to:
  - 1. Run a model and score it
  - 2. Shuffle this feature, re-run the model and score it
    - If the performance significantly dropped, the feature is important and you shoudn't have dropped it
    - If the performance didn't decrease a lot, the feature may be discarded.

#### ? Questions ?

- 1. Perform a feature permutation to detect which features bring the least amount of information to the model.
- 2. Remove the weak features from your dataset until you notice model performance dropping substantially
- 3. Using your new set of strong features, cross-validate a new model, and save its score under variable name strong\_model\_score.

```
import numpy as np
from sklearn.model_selection import cross_validate
from sklearn.inspection import permutation_importance

# Evaluate your model without feature permutation
model = LogisticRegression()
cv_results = cross_validate(model, X, y, cv = 5)
score = cv_results["test_score"].mean()
print(f"Before any feature permutation, the cross-validated accuracy is equal to {
## Question 1 - Permutation importance
model = LogisticRegression().fit(X,y) # Fit the model
```

Before any feature permutation, the cross-validated accuracy is equal to 0.84 After feature permutation, here are the decreases in terms of scores:

Out[50]:

	feature	feature_importance
3	curbweight	0.289267
2	carwidth	0.106073
5	stroke	0.029895
11	enginetype_ohcf	0.018429
6	peakrpm	0.015497
10	enginetype_ohc	0.015445
13	enginetype_rotor	0.011885
0	aspiration	0.008272
4	cylindernumber	0.008168
7	enginetype_dohc	0.004974
12	enginetype_ohcv	0.000419
1	enginelocation	0.000314
8	enginetype_dohcv	0.0
9	enginetype_l	0.0

```
In [51]: ## Question 2 - remove weak features
         # I want to get rid of features which caused less than this in terms of performance
         threshold = 0.05
         # Decompose this one-liner piece of code step by step if you don't understand it at
         weak features = importance df[importance df.feature importance <= threshold]["feature")</pre>
         weak features
         array(['stroke', 'enginetype ohcf', 'peakrpm', 'enginetype ohc',
Out[51]:
                 'enginetype_rotor', 'aspiration', 'cylindernumber',
                 'enginetype_dohc', 'enginetype_ohcv', 'enginelocation',
                 'enginetype_dohcv', 'enginetype_l'], dtype=object)
         ## Question 3 - Cross validating the model with strong features only
In [52]:
         X strong features = df.drop(columns=list(weak features) + ["price"])
         print(f"Our strong features are {list(X strong features.columns)}")
         model = LogisticRegression()
         scores = cross val score(model, X strong features, y, cv = 5)
         strong model score = scores.mean()
         print(f"Before removing weak features, the cross-validated accuracy was equal to {
```

```
print(f"The LogisticRegression fitted with the strong features only has a score of
#### NOTE - The score may even be better because
### some features were bringing nothing else than noise to the model
```

Our strong features are ['carwidth', 'curbweight']
Before removing weak features, the cross-validated accuracy was equal to 0.84
The LogisticRegression fitted with the strong features only has a score of 0.91

#### Test your code

```
In [53]: from nbresult import ChallengeResult
       result = ChallengeResult('strong_model',
                            score = strong_model_score
       result.write()
       print(result.check())
       platform linux -- Python 3.10.6, pytest-7.1.3, pluggy-1.0.0 -- /home/joharlewago
       n/.pyenv/versions/lewagon/bin/python3
       cachedir: .pytest_cache
       rootdir: /home/joharlewagon/code/UKVeteran/05-ML/02-Prepare-the-dataset/data-car-p
       rices/tests
       plugins: anyio-3.6.2, asyncio-0.19.0, typeguard-2.13.3
       asyncio: mode=strict
       collecting ... collected 1 item
       test_strong_model.py::TestStrong_model::test_strong_model_score PASSED
       You can commit your code:
       git add tests/strong_model.pickle
       git commit -m 'Completed strong_model step'
       git push origin master
```

# Bonus - Stratifying your data 🕸

As we split our data into training and testing, we need to be mindful of the proportion of categorical variables in our dataset - whether it's the classes of our target y or a categorical feature in X.

Let's have a look at an example 👇

? Split your original X and y into training and testing data, using sklearn's train\_test\_split; use random\_state=1 and test\_size=0.3 to have comparable results.

```
In [54]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_statest_split)
```

? Check the proportion of price class 1 cars in your training dataset and testing dataset.

If you check the proportion of them in the raw  $\,$  df  $\,$ , it should be very close to  $\,$  50/50

```
In [55]: print('Training data share of class 1 cars:', y_train.mean())
    print('Testing data share of class 1 cars:', y_test.mean())

Training data share of class 1 cars: 0.5037593984962406
    Testing data share of class 1 cars: 0.5172413793103449
```

It should still be pretty close to 50/50 🤞

#### But what if we change the random state?

? Loop through random states 1 through 10, each time calculating the share of price class 1 cars in the training and testing data. ?

```
In [56]: for i in range(1, 11):
    print("-"*50)
    print("#### Random state set =", i)
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random print('Training data share of class 1 cars:', round(y_train.mean(), 3))
    print('Testing data share of class 1 cars:', round(y_test.mean(), 3))
```

Training data share of class 1 cars: 0.504

##### Random state set = 1

```
Testing data share of class 1 cars: 0.517
        ______
        ##### Random state set = 2
        Training data share of class 1 cars: 0.481
        Testing data share of class 1 cars: 0.569
        ______
        ##### Random state set = 3
        Training data share of class 1 cars: 0.504
        Testing data share of class 1 cars: 0.517
        ##### Random state set = 4
        Training data share of class 1 cars: 0.534
        Testing data share of class 1 cars: 0.448
         ##### Random state set = 5
        Training data share of class 1 cars: 0.534
        Testing data share of class 1 cars: 0.448
         ##### Random state set = 6
        Training data share of class 1 cars: 0.496
        Testing data share of class 1 cars: 0.534
        ##### Random state set = 7
        Training data share of class 1 cars: 0.534
        Testing data share of class 1 cars: 0.448
        ##### Random state set = 8
        Training data share of class 1 cars: 0.489
        Testing data share of class 1 cars: 0.552
        _____
        ##### Random state set = 9
        Training data share of class 1 cars: 0.579
        Testing data share of class 1 cars: 0.345
        ##### Random state set = 10
        Training data share of class 1 cars: 0.489
        Testing data share of class 1 cars: 0.552
        You will observe that the proportion changes every time, sometimes even quite drastically
        ! This can affect model performance!
         ? Compare the test score of a logistic regression when trained using
        train_test_split(random_state=1) vs. random_state=9 ?
        Remember to fit on training data and score on testing data.
In [57]: model 1 = LogisticRegression()
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_st
        model_1.fit(X_train, y_train)
        model_1.score(X_test, y_test)
        0.9310344827586207
Out[57]:
In [58]:
        model_9 = LogisticRegression()
        X train, X test, y train, y test = train test split(X, y, test size=0.3, random sta
```

```
model_9.fit(X_train, y_train)
model_9.score(X_test, y_test)
```

Out[58]:

0.7931034482758621

You should see a much lower score with random\_state=9 because the proportion of class 1 cars in that test set is 34.5%, quite far from the 57.9% in the training set or even the 50% in the original dataset.

This is substantial, as this accidental imbalance in our dataset can not only make model performance worse, but also distort the "reality" during training or scoring 

©

So how do we fix this issue? How do we keep the same distribution of classes across the train set and the test set?

Luckily, this is taken care of by cross\_validate in sklearn, when the estimator (a.k.a the model) is a classifier and the target is a class. Check out the documentation of the cv parameter in sklearn.model\_selection.cross\_validate.

The answer is to use the following:

Stratification

### Stratification of the target

- We can also use the **strafification** technique in a train\_test\_split.
- ? Run through the same 1 to 10 random state loop again, but this time also **pass** stratify=y into the holdout method. ?

```
In [59]: for i in range(1, 11):
    print("-"*50)
    print("#### Random state set =", i)
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random
    print('Training data share of class 1 cars:', round(y_train.mean(), 3))
    print('Testing data share of class 1 cars:', round(y_test.mean(), 3))
```

```
##### Random state set = 1
Training data share of class 1 cars: 0.511
Testing data share of class 1 cars: 0.5
_____
##### Random state set = 2
Training data share of class 1 cars: 0.511
Testing data share of class 1 cars: 0.5
______
##### Random state set = 3
Training data share of class 1 cars: 0.511
Testing data share of class 1 cars: 0.5
______
##### Random state set = 4
Training data share of class 1 cars: 0.511
Testing data share of class 1 cars: 0.5
______
##### Random state set = 5
Training data share of class 1 cars: 0.511
Testing data share of class 1 cars: 0.5
##### Random state set = 6
Training data share of class 1 cars: 0.511
Testing data share of class 1 cars: 0.5
##### Random state set = 7
Training data share of class 1 cars: 0.511
Testing data share of class 1 cars: 0.5
-----
##### Random state set = 8
Training data share of class 1 cars: 0.511
Testing data share of class 1 cars: 0.5
_____
##### Random state set = 9
Training data share of class 1 cars: 0.511
Testing data share of class 1 cars: 0.5
##### Random state set = 10
Training data share of class 1 cars: 0.511
Testing data share of class 1 cars: 0.5
```

• Even if the random state is changing, the proportion of classes inside the training and testing data is kept the same as in the original y. This is what *stratification* is.

Using train\_test\_split with the stratify parameter, we can also preserve proportions of a feature across training and testing data. This can be extremely important, for example:

- preserving proportion of male and female customers in predicting churn
- preserving distribution of 1-5 review scores (multiclass!) in recommending the next product
- etc...

For instance, in our dataset, to holdout the same share of aspiration feature in both training and testing data, we could simply write train\_test\_split(X, y, test\_size=0.3, stratify=X.aspiration)

We need StratifiedKFold 🔬

## Stratification - generalized

StratifiedKFold allows us to split the data into K splits, while stratifying on certain columns (features or target).

This way, we can do a manual cross-validation while keeping proportions on the categorical features of interest - let's try it with the binary aspiration feature:

```
In [60]: from sklearn.model_selection import StratifiedKFold

# initializing a stratified k-fold that will split the data into 5 folds
skf = StratifiedKFold(n_splits=5)
scores = []

# .split() method creates an iterator; 'X.aspiration' is the feature that we strate
for train_indices, test_indices in skf.split(X, X.aspiration):

# 'train_indices' and 'test_indices' are lists of indices that produce proporte
X_train, X_test = X.iloc[train_indices], X.iloc[test_indices]
y_train, y_test = y.iloc[train_indices], y.iloc[test_indices]

# initialize and fit a model
model = LogisticRegression()
model.fit(X_train, y_train)

# append a score to get an average of 5 folds in the end
scores.append(model.score(X_test, y_test))

np.array(scores).mean()
```

Out[60]:

- 0.8585695006747638
- Some sklearn reads on **stratification**:
  - Visualization of how different holdout methods in sklearn work
  - Overall cross-validation and stratification understanding
- **XXX** Congratulations! You have prepared a whole dataset, ran feature selection and even learned about stratification &
- Place Don't forget to git add/commit/push your notebook...
- ... and move on to the next challenge!